



POLITECNICO
MILANO 1863

SCUOLA DI INGEGNERIA INDUSTRIALE
E DELL'INFORMAZIONE

TESI DI LAUREA MAGISTRALE IN
COMPUTER SCIENCE AND ENGINEERING
INGEGNERIA INFORMATICA

Binary Instrumentation for Tracking Edge Coverage in Fuzzing

Author:

Scandroglio Edoardo

Student ID:

216600

Advisor:

Prof. Name Surname

Academic Year:

2026-27

Dedicated to someone.

Abstract

...

Keywords: key, words, go, here

Abstract in lingua italiana

...

Parole chiave: qui, vanno, le, parole, chiave

Contents

Abstract	i
Abstract in lingua italiana	iii
Contents	v
1 Background and Motivation	1
1.1 Fuzzing Background	1
1.2 Binary Rewriting Background	4
1.3 ELF Loading and Relocations	7
1.3.1 ELF Structure	7
1.3.2 ELF Linking	8
1.3.3 ELF Loading	9
1.3.4 ELF Symbols and Relocations	11
1.4 Problem Statement	14
1.5 State of the Art	15
1.5.1 AFL Functioning	17
1.6 Goals and Challenges	19
2 Approach	21
2.1 Forkserver Insertion	21
2.2 Shared Memory Instrumentation	24
2.3 Instrumentation Points Search	26
2.4 Actual Instrumentation	30
2.5 Forkserver Optimization	32
3 Implementation	35
3.1 Forkserver Specification	35
3.2 Instrumentator Specification	37

3.2.1	Binary Modification	38
3.2.2	Instrumentation	40
4	Experiments	47
4.1	Goals	47
4.2	Datasets and Experiment Setup	48
4.3	Experiment I	49
4.4	Experiment II	50
5	Limitations and Future Works	53
5.1	Limitations	53
5.2	Future Works	54
	Bibliography	55
	List of Acronyms	59
	List of Figures	61
	List of Tables	63
	List of Algorithms	65

1 | Background and Motivation

This chapter will focus on presenting the context of this thesis. In particular, in Section 1.1 the fundamentals of fuzzing will be discussed, with particular attention to coverage-guided grey-box fuzzing. In Section 1.2 the background on binary rewriting is discussed, along with its approaches and difficulties. In the next chapter will be presented the approach used in order to resolve the main problems regarding the solution discussed in this thesis. In order to do so some concept of binary loading must be presented (Section 1.3). Section 1.4 states the problem addressed by this work, and Section 1.5 reviews the state of the art in binary-only fuzzing, also focusing on a description of AFL/AFL++ internals. In the last section (Section 1.6) I will present the main goals and challenges of this thesis, and explain why the implemented tool differs from the state of the art.

1.1. Fuzzing Background

Fuzzing is nowadays a widely used approach for software testing and vulnerability assessment [19, 31]. Its main strength derives from the fact that it can rapidly produce various inputs to be fed into a target executable. In general, a fuzzer can be divided into the **Input Generator**, which is responsible for providing inputs to the **Executor** that executes them. During the execution, the **Defect Monitor** checks whether the executed input causes a crash of the target.

In particular, the Input Generator must choose which input to provide to the executor among all the inputs in the *Input Space*. A subset of the *Input Space* is represented by the *Valid Input Space*. This is a subset of the input space that specifically contains all the inputs that are well formed with respect to the application being tested. For example, if the fuzzed target is a PDF reader, a valid input is a PDF file. A good Input Generator must choose inputs that are also valid inputs, because only by doing this is there a probability of finding a vulnerability [31]. Hence the need for a precise input selection and mutation strategy arises, in order to discover as many edge cases as possible.

In order to do so, the Input Generator must obtain some information from the Executor. Depending on the amount of information the Input Generator has, fuzzing can be classi-

fied into *White-box*, *Black-box* or *Grey-box*. In Black-box fuzzing, the fuzzer does not have any knowledge of the execution [21]. In White-box fuzzing, instead, the source code must be available in the general case, and techniques such as Symbolic and Concolic execution are used [13]. In Grey-box fuzzing, instead, the inputs are chosen using, for example, **Code Coverage** information (i.e. the number of lines of code visited during execution of the target). This approach is called *Coverage-Guided Grey-box fuzzing* and is based on prioritizing the inputs that discover new sections of code. Since the source code is not always available, Code Coverage is not always a viable Coverage Measure. In general, the coverage measures can be of different granularity [12]. The three main options are:

- **Block Coverage:** Keeps track of the hit count for each Basic Block (BB). This approach is used by some popular fuzzers such as libFuzzer or honggfuzz [12]. This method is often too imprecise because it fails to capture the transitions between BBs.
- **Edge Coverage:** Keeps track of the hit count of every edge in the Control Flow Graph (CFG) (i.e. the transition between two BBs). This is the approach used by AFL and AFL++ [10]. It maintains overheads similar to Block Coverage but can track the transitions between BBs and not only their execution, adding some context information. However, it lacks precision in keeping track of the order in which the edges are taken.
- **Path Coverage:** This is the most complete metric and keeps track of the complete path taken during the execution of a specific input, thus also keeping track of the order in which the edges are taken. The main drawback of this metric is that a single execution path could become very large, resulting in large runtime and memory overheads.


One of the best approach to Code Coverage, in terms of tradeoff between overhead and effectiveness, is **Edge Coverage** [12], which keeps track of which edges in the CFG are taken during execution without the overheads incurred by keeping track of the entire execution path. Coverage information is obtained through the instrumentation of the target application. Depending on the availability of the source code, the instrumentation can be added in different ways:

- The source code is available: In this case the instrumentation is added at *Compile Time*.
- The source code is not available: In this case there are two different techniques for instrumenting the application:

- **Static Binary Instrumentation (SBI):** The instrumentation is added before the target is executed, without access to the source code, relying on Static Binary Rewriting techniques.
- **Dynamic Binary Instrumentation (DBI):** The instrumentation is added during program execution, relying on Dynamic Binary Rewriting techniques.

Regardless of the availability of the source code, the *Coverage-Guided Grey-box fuzzing* technique can be summarized as shown in fig. 1.1.

Figure 1.1: Generic diagram describing the Coverage-Guided Grey-box fuzzing workflow



Specifically for **Edge Coverage**, fig. 1.1 can be described as follows:

1. **Seed Corpus:** Represents the set of inputs on which mutations are computed. At the beginning it is composed of user-defined seeds, but during fuzzing this set can grow if a specific mutation discovers new edges not yet found.
2. **Input Generator:** Composed of two main parts:
 - **Seed Scheduling:** responsible for deciding which seed in the corpus to mutate.
 - **Mutation Scheduling:** responsible for deciding which mutation or mutations to apply to the selected seed.

3. **Executor:** The component that actually executes the target binary. If the target binary has already been instrumented with Edge Coverage tracking instructions, it is executed as is; otherwise, in the case of DBI, a dynamic translator must dynamically instrument the binary.
4. **Coverage Collection:** In the case of Edge Coverage, this part is represented by the data structure used to keep track of which edges between BBs have been taken and which have not.
This information is used, in some tools such as AFL++ [10], to guide Seed Scheduling, a technique that, given different inputs with the same level of coverage (i.e. number of edges visited), prioritizes the input that is smaller in size and takes less execution time.
5. **Defect Monitor:** This component checks the exit status of the target to catch any crash or timeout of the target, or any other unexpected behavior. If none of these conditions are met, meaning that the target ends normally, the coverage is checked to see whether the current input discovered a new edge. If the input discovered a new edge, that input is added to the corpus.

1.2. Binary Rewriting Background


Binary rewriting consists in altering an already compiled executable without having access to the source code, while ensuring that the modified binary remains executable. The process of binary rewriting can be classified into Static Binary Rewriting, if the binary is stored in persistent memory while being rewritten, and Dynamic Binary Rewriting, if the binary is being executed. The process of binary rewriting can be divided into four main steps [27]:

1. *Parsing:* In this phase the main goal is to retrieve information about the binary. In particular, the output of this phase is a raw stream of instructions (in machine code) and data. This information is then used in the Analysis step.
2. *Analysis:* The information gained from the Parsing phase is analyzed to reconstruct the program structure. For example, the raw stream of machine code is passed through a disassembler that reconstructs the assembly instructions. A structural recovery is performed, which includes function detection and CFG reconstruction, and labels and symbols are extracted from the raw stream of data.
3. *Transformation:* Once the CFG has been reconstructed and function and data boundaries have been recovered, the binary rewriter can modify the binary at the

Instrumentation Point (IP). These are user-defined locations in the binary where the rewriter must add, modify or delete an instruction. The IP can be coarse- or fine-grained. In the first case, for example, the instrumentation could be inserted at branch level (at the beginning of each BB). In the second case, instead, any instruction could potentially be modified.

4. *Code Generation*: In this last step, the changes made in the previous steps must be integrated into the binary while preserving its executability. This can be done in different ways. For example, the changes could be inserted into a new section of the executable, redirecting execution to the new code at the IP. Alternatively, this step could also be performed by reassembling a completely new executable starting from the changes made at the IP. This second approach is much harder, because it requires preserving all the references inside the executable and preserving the executable layout.

Figure 1.2: An example of a Static Binary Rewriting Workflow




In fig. 1.2, a general scheme for static binary rewriting can be seen. In this case all

the phases occur prior to execution, so the cost in terms of overhead is paid only once, before the execution. The main problems are disambiguating between code and data (i.e. deciding whether a stream of bytes represents machine code instructions or constants placed inline in the .text section) and resolving indirect branches [3, 27]. This problem is generally undecidable and can only be solved using heuristics. There are three main families of static transformation:

- **Direct:** Consists in directly modifying the instruction. Works well with RISC systems with instructions having the same length. Generally fails when adding or removing instructions, because this does not preserve the offset into the .text section.
- **Minimal-invasive:** The new code is inserted into a new section of the binary, and at the IP an unconditional jump to the new section is placed. The main limitation of this approach arises when the new section is far from the IP and a near jump (necessary to reach the new section) uses a different number of bytes than the instruction at the instrumentation point. In fact, in x86-64 assembly a short jump uses 2 bytes while a near jump uses 5 bytes. So, if a short jump is needed at the instrumentation point for alignment reasons, but the new section is too far away and a near jump is needed instead, it is necessary to add a **trampoline**. This approach can be seen in fig. 1.3. Trampoline-based instrumentation is one of the most common approaches to binary rewriting [27].
- **Full-translation:** The binary is lifted into an Intermediate Representation (IR) that makes it possible to add an IP at instruction level. The IR is then reassembled into machine code after the changes are made.

Figure 1.3: An example of Static Instrumentation using the trampoline technique



With regard to Dynamic Binary Rewriting, the Analysis and Transformation phases are

executed at runtime for each instruction. In this case, none of the problems that arise in Static Binary Rewriting occur, since the binary is being executed, so indirect branches can be resolved using the values in the registers, and there is no need to disambiguate between constants inside the `.text` section and instructions. The main drawback of this approach is the prohibitive overhead caused by the dynamic analysis and transformation of each instruction, together with the fact that only the taken branches can be modified. In fact, if a branch is not actually taken during execution, its instructions never reach the rewriter.

A specific type of binary rewriting is Binary Instrumentation. Binary instrumentation uses binary rewriting techniques to add instructions that track binary behaviour, log events, or, in the case of this thesis, insert edge coverage tracing instructions.

1.3. ELF Loading and Relocations

Executable and Linkable Format (ELF) is the standard executable format under Linux operating systems. The work done in this thesis heavily involves the ELF structure, with particular focus on some sections used for relocation, alongside the process of executing an ELF. For this reason the following section will disclose in details some of the ELF loading and relocation mechanism. But, in order to do so, a review of the ELF basic structure must be done.

1.3.1. ELF Structure

In this subsection the process of ELF execution is discussed, with a focus on the parts of the execution flow that will influence this work.

The **ELF Header** is a structure located at the beginning of the ELF and contains some important information about the type of ELF, the target operating system Alongside this information there is the location of two important data structures inside the ELF:

- The **Program Header Table**, which defines the *Execution View*
- The **Section Header Table**, which defines the *Linking View*

The **Program Header Table** contains information needed to create a runtime image of the program in memory, via the **Segments** (*i.e. which parts of the file must be loaded into memory, which shared objects are needed . . .*). Each segment contains the access flags for that part of memory (i.e. Read/Write/Execute). The **Section Header Table** contains information about the **Sections** that compose the binary, together with their mapping to the segments (e.g. which segment the `.text` section is mapped to). Among all

the sections, some contain specific structures used to perform specific tasks:

- **.dynamic**: The dynamic section contains information used by the Dynamic Linker (e.g. the dependencies, the interpreter location ...).
- **.dynsym**: List of symbols defined and used by the binary. If the binary is not stripped, there is also the **.symtab** section, which contains a more complete list of symbols.
- **.hash** or **.gnu.hash**: Structures used for efficient symbol lookup during symbol resolution¹.
- **.rela.dyn** and **.rela.plt**: Both are relocation tables. The first is the standard one, the second is used in the case of **Lazy Binding**.
- **.shstrtab**, **.dynstr** and **.strtab**: Hold lists of strings referenced by other parts of the ELF. The first holds section names, the second holds the names of the symbols contained in the **.dynsym** section, and the last one holds the names of the entire **.symtab** section.
- **.got** and **.got.plt**: Contain the Global Offset Table (GOT), filled by the dynamic linker with the addresses of external functions and variables. The first is the standard one, the second is used by the Procedure Linkage Table (PLT).
- **.plt**: Contains the PLT, which is an executable section made up of trampolines. In particular, it contains stubs for functions that might be located in another ELF or Shared Object (SO). The stub references the corresponding GOT entry containing the resolved address of the function. This section is used in the case of **Lazy Binding**.
- **.init** and **.init_array**: Sections used for initialization. **.init** contains initialization code, while **.init_array** contains the addresses of initialization functions.
- **.fini** and **.fini_array**: Sections used for termination handling. **.fini** contains termination code, while **.fini_array** contains the addresses of termination functions.

1.3.2. ELF Linking

Before actually digging deep into ELF loading, the linking process must be reviewed. The **Linking phase** consists of joining together different pieces of code belonging to different files. For example, if a program uses the **printf** function, it needs to take the code of **printf** from the **libc** library. This process of joining together the executable that

¹The **.gnu.hash** section is the more efficient version of **.hash**.

uses `printf` with the `printf` code located in another file is called **Linking**. Linking can be of two different types:

- **Static linking**: The needed code (e.g. the instructions of `printf`) is copied inside the binary at *Link Time*. This causes the size of the binary to be bigger (because the code of external functions actually lies inside it) and consequently wastes RAM space (if 10 binaries all use the `printf` function, the code of `printf` is copied inside all 10 binaries).
- **Dynamic linking**: The needed code is not copied inside the binary; instead, the binary contains information (i.e. the address of the external function or variable) used to locate the needed resource. The resolution happens at runtime, when the Dynamic Linker finds and loads the library into the process's virtual address space and resolves the references to the symbols. This allows different processes to share the same SO.


Since nowadays the most commonly used approach is **Dynamic Linking**, the following discussion will focus on it.

As stated before, the linking process can be used for both external variables and external functions. In both cases the process is the same, since both functions and variables can be exported as symbols.


When dealing with functions there is the possibility of using **Lazy Binding**. This technique consists of resolving the symbol only during the first invocation of the function, rather than at load time as happens with variables. In order to achieve **Lazy Binding** a stub section is used (i.e. `.plt`). When a program uses a library function, a PLT entry is added to the `.plt` section. As can be seen in fig. 1.4, the relocation is resolved lazily upon the first invocation of the function. When control is transferred to the corresponding PLT entry, the entry performs an indirect jump through its associated GOT entry. Before the symbol is resolved, this GOT entry contains an address that redirects execution back to the PLT, eventually reaching the PLT resolver stub. The resolver then invokes the dynamic linker, which performs the symbol lookup, resolves the relocation, and updates the GOT entry with the resolved address of the function, as can be seen in fig. 1.4b. Subsequent calls can therefore jump directly to the resolved address through the GOT.

1.3.3. ELF Loading

The next useful step in order to understand the chosen approach consists of going into detail on the ELF Loading phase. There are 3 main actors that allow a binary to be executed properly:



(a) An example of puts function resolution using Dynamic Linking and Lazy Binding before the actual resolution



(b) Same binary as fig. 1.4a but with symbol resolved

Figure 1.4: An example of Lazy Binding involving .plt and .got

- **Kernel:** responsible for the actual `execve` invocation and allocation of the different memory segments. Moreover, it is the one that executes the dynamic linker used to resolve dependencies prior to the execution of the binary.
- **Dynamic Linker:** Loads all the needed libraries into the process memory, resolves the external dependencies, prepares the .got and .plt, and executes the initializers of the libraries.
- **Target Binary:** Invoked by the Dynamic Linker, it prepares the environment for the main function and, in the end, runs it.

All these procedures can be seen in detail in fig. 1.5. The dynamic linker part is the one on which I will focus in this dissertation. The dynamic linker:




Figure 1.5: Dynamically linked ELF loading procedure

- *Loads all the DT_NEEDED entries:* these are the dependencies (i.e. the SOs needed by the binary). For example, if the binary uses an external function contained in *libc.so*, the library from which the function is exported must be added as a dependency.
- *Resolves the relocations:* a focus on relocations will be presented in the next sub-section.
- *Arranges the GOT and PLT.*
- *Executes all the initializer functions (constructors) for each dependency:* For each *DT_NEEDED* entry found by the dynamic linker, it executes all the functions contained in the *.init_array* section of that entry. These are functions declared with `__attribute__((constructor))`.
- *Finds and executes the binary entry point.*

1.3.4. ELF Symbols and Relocations

The dynamic linker uses symbols to link different ELF's together.

- A component (an ELF or a SO) can export a symbol representing a variable or a function inside it.
- Another component can reference the exported symbol and use it (if the first component is a dependency of the second).
- The linker acts as the intermediary between the two components, responsible for matching the symbol names (the exported one and the referenced one) and storing the address of the exported symbol in the second component's address space in such a way that this component can access it.

All the symbols belonging to an ELF (both imported and exported) are stored in the *.dynsym* section. In fig. 1.6 a list of symbols extracted from a SO can be seen. The

20:	0000000000000000	0	FUNC	GLOBAL	DEFAULT	UND	exit@GLIBC_2.2.5 (2)
21:	0000000000000000	0	NOTYPE	WEAK	DEFAULT	UND	_ITM_registerTMC[...]
22:	0000000000000000	0	FUNC	WEAK	DEFAULT	UND	[...]@GLIBC_2.2.5 (2)
23:	0000000000000000	0	FUNC	GLOBAL	DEFAULT	UND	fork@GLIBC_2.2.5 (2)
24:	00000000000040b8	4	OBJECT	GLOBAL	DEFAULT		25 standalone
25:	00000000000013f2	373	FUNC	GLOBAL	DEFAULT		14 aflplusplus_fork[...]
26:	0000000000001891	137	FUNC	GLOBAL	DEFAULT		14 afl_forkserver_c[...]
27:	00000000000016fd	404	FUNC	GLOBAL	DEFAULT		14 init
28:	00000000000040b0	8	OBJECT	GLOBAL	DEFAULT		25 __forkserver_sta[...]
29:	0000000000001567	252	FUNC	GLOBAL	DEFAULT		14 afl_forkserver_start
30:	0000000000001663	154	FUNC	GLOBAL	DEFAULT		14 standalone_forkserver
31:	00000000000040a8	8	OBJECT	GLOBAL	DEFAULT		25 __afl_area_ptr

Figure 1.6: Example of a *.dynsym* table for a shared object

Relocation section '.rela.dyn' at offset 0x7ddaac contains 10 entries:				
Offset	Info	Type	Sym. Value	Sym. Name + Addend
000000808ff8	002300000006	R_X86_64_GLOB_DAT	0000000000000000	__gmon_start__ + 0
000000809280	004000000005	R_X86_64_COPY	0000000000809280	__progname@GLIBC_2.2.5 + 0
000000809288	003f00000005	R_X86_64_COPY	0000000000809288	stdout@GLIBC_2.2.5 + 0
000000809290	004100000005	R_X86_64_COPY	0000000000809290	stdin@GLIBC_2.2.5 + 0
000000809298	004600000005	R_X86_64_COPY	0000000000809298	optind@GLIBC_2.2.5 + 0
0000008092a0	004700000005	R_X86_64_COPY	00000000008092a0	optarg@GLIBC_2.2.5 + 0
0000008092a8	004300000005	R_X86_64_COPY	00000000008092a8	__progname_full@GLIBC_2.2.5 + 0
0000008092c0	004500000005	R_X86_64_COPY	00000000008092c0	stderr@GLIBC_2.2.5 + 0
0000017db000	003d00000006	R_X86_64_GLOB_DAT	0000000000000000	__afl_area_ptr + 0
0000017db008	003e00000006	R_X86_64_GLOB_DAT	0000000000000000	__forkserver_start_ptr + 0

Figure 1.7: Example of a *.rela.dyn* relocation table extracted from a binary

first 4 symbols are external symbols that must be resolved. This can be seen because the Value column is 0 and Ndx is UND (undefined). The last 8, instead, are symbols declared inside the shared object. When the visibility is set to *DEFAULT*, this means that the symbol can be visible outside the shared object.

The process of resolving symbol references is called **Relocation**. The relocation algorithm is shown in algorithm 1.1. Relocating a symbol means resolving the symbol (i.e. finding the address of that symbol in the original object), and writing that address inside the binary that has the relocation. Specifically, the symbol address is written inside the GOT. In fig. 1.7, the address at which the symbol address must be written can be seen in the “Offset” column.

Algorithm 1.1 Relocation resolution at ELF load time

Require: $L = [obj_1, \dots, obj_n]$ {The objects that have a dependency (DT_NEEDED) must already be loaded in memory}

Ensure: for each $obj \in L$, the respective GOT entry contains the resolved runtime address of the external symbol.

```

1: for  $obj \in L$  do
2:    $Reloc\_table \leftarrow obj.GET\_SECTION(".rela.dyn")$ 
3:   for  $reloc \in Reloc\_table$  do
4:     if  $reloc.type = R\_*\_RELATIVE$  then
5:        $addr \leftarrow obj.loadBias + r.addend$  {The symbol does not need to be resolved.
        The address must be adjusted in order to enable ASLR}
6:     else
7:       if  $reloc.type \in \{R\_*\_GLOB\_DAT, R\_*\_JUMP\_SLOT\}$  then
8:          $sIdx \leftarrow reloc.symbol\_index$ 
9:          $symbolName \leftarrow obj.dynstr[obj.dynsym[sIdx].name]$  {The name of the
          symbol is read from .dynsym}
10:         $target \leftarrow RESOLVESYMBOL(symbolName, L, obj)$ 
11:        if  $target = NULL$  then
12:          error: undefined symbol  $symbolName$ 
13:        else
14:           $(defObj, defIdx) \leftarrow target$  {defObj is the object that exports the symbol,
          defIdx is the index in the .dynsym table of defObj}
15:           $addr \leftarrow defObj.loadBias + defObj.dynsym[defIdx].value$  {The address of the
          resolved symbol is computed as base_address (of the object that exports the
          symbol) plus the offset at which the symbol is stored}
16:        end if
17:      end if
18:    end if
19:    write  $addr$  in  $obj.got$  at offset  $reloc.offset$  {The destination address at which to
      write the resolved symbol address is specified in the relocation entry}
20:  end for
21: end for

```

Algorithm 1.2 RESOLVESYMBOL function used by algorithm 1.1

Require: name *name*, object list *L*, object requesting the resolution *requester*
Ensure: pair (*obj*, *idx*) of the first symbol found, or NULL if the symbol has not been defined by any object.

```

1: for obj ∈ L {The order is: first the executable, then all the DT_NEEDED dependencies
   in order of definition} do
2:   idx ← GNULIBLOOKUP(name, obj.gnuHash, obj.dynsym, obj.dynstr) {As seen before,
   the hashes are used in order to speed up the symbol lookup}
3:   if idx ≠ NOT_FOUND && obj.dynsym[idx].st_shndx ≠ SHN_UNDEF then
4:     return (obj, idx)
5:   end if
6: end for
7: return NULL

```

1.4. Problem Statement

Coverage-Guided Grey-box fuzzing is one of the most effective techniques for automatic vulnerability discovery in software [6, 22], due to its ability to generate a wide range of inputs based on the feedback obtained through Code Coverage information, which guides the fuzzer towards unexplored regions of code and therefore helps it find hidden vulnerabilities. In fact, *Black-box* fuzzing lacks a precise mutation strategy based on execution monitoring, while *White-box* fuzzing incurs high overheads due to Symbolic or Concolic execution. Furthermore, the latter approach also requires the availability of the source code in the general case, which is not always available. So *Grey-box* fuzzing seems to be the most promising trade-off between precision of input mutation and overhead. However, *Grey-box* fuzzing also has its own drawbacks. As a matter of fact, the effectiveness of this approach strongly depends on the availability of the source code. In fact, both SBI and DBI have many limitations. The main limitation of DBI is the huge amount of overhead (higher than 600% in the case of AFL++ in QEMU mode [22]), which makes the exploration of a substantial part of the valid input space almost impossible. Instead, the main difficulty of SBI arises in modifying already compiled binaries. There are several problems regarding Binary Rewriting, such as Indirect Control Transfers (e.g. `call rax`), shifts in binary sections, or instructions that require rewriting the instructions that reference these sections, and many others, as seen in the previous Section 1.2.

For all these reasons, the preferred technique nowadays is *Compile-time Edge Coverage instrumentation* with fuzzers such as AFL++, LibAFL [11], or libFuzzer [17]. However, in practice, for reasons of industrial secrecy (i.e. proprietary software) or due to precompiled libraries, applications are often closed source, and so Compile-time Edge Coverage instrumentation becomes unusable.

In these scenarios, binary-only fuzzing is the only viable approach. Although SBI presents many challenges, it nonetheless proves to be the most fruitful area, given its low overheads [22]. In this thesis I will address the problem of statically instrumenting a binary with Edge Coverage instructions in a way that makes the instrumented binary compatible with AFL++, which has been chosen as the fuzzer due to its popularity and because it represents the current state of the art, implementing most of the recently discovered techniques demonstrated by Andrea Fioraldi et al. [10].

The instrumentation must add as little overhead as possible with respect to compile-time overheads and must preserve the correctness of the binary. This is important in fuzzing, because lower overhead means more executions per second, so a larger portion of the valid input space can be explored.

1.5. State of the Art

In this section I will present some of the current state-of-the-art solutions in binary-only grey-box fuzzing, distinguishing between SBI and DBI approaches.

- **DBI**

- **AFL++ in QEMU/FRIDA mode** [10]: AFL++ in QEMU mode uses QEMU in *User-mode Emulation*, which consists in translating guest machine code (i.e. guest BBs) into an intermediate representation called Translation Block (TB) and subsequently appending the coverage tracing instructions at the end of it. This approach is based on emulating the entire CPU using QEMU’s virtualized environment. As stated before, AFL in QEMU mode has overheads on the order of 600% with respect to compile-time instrumentation [22].

AFL++ in FRIDA mode [1], instead, is based on shared library injection, in which a Frida [23] component called **Stalker** intercepts any BB that is about to be executed, copies it into a **Code Cache** and dynamically instruments the copy with coverage tracing instructions. Execution control is then transferred to the copy, which is executed and then re-intercepted by **Stalker** at the end. Unlike QEMU mode, it does not emulate the CPU, but instead executes within the memory space of the target program. AFL++ reports that FRIDA mode is approximately 2 to 5 times slower than compile-time instrumentation [2].

- **DynamoRIO/Pin** [4, 18]: Both these approaches share the same basic principle, called **Just in-time (JIT) code caching**. This consists in taking control of the thread when the target application starts, and from that moment on,

every BB is intercepted before it is executed and is copied into a private **Code Cache**. The instrumentation is then added to the copy, and control is transferred to the code cache. The instrumentation process is more or less identical to FRIDA’s Stalker mechanism, but is not as optimized. Also in this case, the implemented version uses AFL [30] as the fuzzer [15, 16].

- **SBI**

- **Retrowrite** [8]: This tool performs symbolic lifting on the code, replacing every reference discovered during disassembly with a symbolic label. Doing this allows instructions to be added without modifying the references, achieving reassembleable assembly. While RetroWrite can add ASan instrumentation, it cannot natively add AFL-style edge coverage instructions. To do so, the lifted assembly must be instrumented with the afl-gcc tool in order to produce the patched binary. Moreover, community-reported issues [20] document missing instrumentation points when RetroWrite’s symbolized output is compiled through afl-gcc, caused by a mismatch between RetroWrite’s label-naming convention and afl-as’s basic-block detection heuristics. Finally, RetroWrite only works with PIE, non-stripped binaries.
- **AFL-Dyninst** [14]: AFL-Dyninst employs Dyninst’s binary rewriting capabilities to instrument an existing executable by inserting callbacks at selected basic blocks. Dyninst’s instrumentation mechanism is based on code patching and trampoline redirection [5], which causes many context switches and state saving/restoring operations. Its performance is slightly better than DBI, but this approach still reaches 500% fuzzing overhead, as demonstrated by Stefan Nagy et al. [22].
- **ZAFL** [22]: This tool uses two main phases: the first consists in reconstructing an IR using the Zipr binary rewriter. The resulting IR is then used by ZAX to perform an optimization on the IR, find the instrumentation points, and finally instrument these points. ZAFL works by inserting calls to a function that updates the shared map at the beginning of each BB. Despite ZAFL’s many improvements, Jifan Xiao et al. [28] demonstrated that it still causes crashes in 28% of the binaries used in their cited paper.

Before proceeding to the next section, in which I will discuss the main goals of this thesis, I want to discuss the core functioning of AFL and AFL++, given its importance for understanding my approach.

1.5.1. AFL Functioning

[29] As previously stated, AFL and, in particular, AFL++ are among the most widely used grey-box fuzzing frameworks [10]. AFL++ represents the community-based continuation of the original AFL project. The core functioning is virtually the same and consists of a Shared Memory (SHM) region shared between the fuzzer and the target. The SHM size can vary, but the default is 65536 bytes.

Each “cell” of the SHM contains the number of times a given edge in the CFG has been taken during the execution of the target. Since the SHM has 65536 “cells”, each “cell” can contain a number from 0 to 255. The “cell” position is computed with a specific hash function applied to the IDs of the two BBs that form the edge. The main Edge Tracing update rule is summarized in algorithm 1.3. Line 3 of algorithm 1.3 is used to avoid

Algorithm 1.3 AFL-style SHM update rule

```

1:  $idx \leftarrow current\_bb\_id \oplus previous\_bb\_id$ 
2:  $shared\_map[idx] \leftarrow shared\_map[idx] + 1$ 
3:  $previous\_bb\_id \leftarrow current\_bb\_id \gg 1$ 

```

updating the same SHM location when an edge is traversed forward or backward. An example of this could be a simple **do-while** loop, as seen in algorithm 1.4 and fig. 1.8.

The main idea behind AFL is to instrument each basic block of a target application with

Algorithm 1.4 Do-While Example

```

repeat
  {Basic Block A}
until Permanence_Condition
{Basic Block C}

```

Without 3 of algorithm 1.3, the hash of Edge1 would result in $ID_A \oplus ID_B$, and the hash of Edge2 (if the permanence condition is not met) would result in $ID_B \oplus ID_A = ID_A \oplus ID_B$, due to the symmetry of XOR. But since the hash is the same, there would be a hash collision that would cause the same “cell” to be updated twice instead of once.




Figure 1.8: Do-while CFG

the shared map update rule. The fuzzer initializes the SHM memory with zeros, runs the

instrumented target, which updates the SHM according to algorithm 1.3 until it reaches the end, and then returns to the fuzzer. The SHM ID is communicated to the target through an environment variable called `__AFL_SHM_ID`.

At this point the fuzzer checks the SHM it has just obtained, called `trace_bits`, searching for any newly discovered edge. This is done by keeping a structure called `virgin_bits`, which contains all the edges not yet discovered and is initialized with ones. The comparison is performed through `trace_bits & virgin_bits`. If any position results in one after the logical AND, this means a new edge has been discovered. In this case, the `virgin_bits` data structure is updated and the input just executed is saved in the corpus. Otherwise, a new input is chosen, mutated and then executed.

This approach requires a `fork()` call followed by an `execve()` call for each input, causing significant overhead for each input execution. Two main optimizations have been implemented:

- **forkserver**
- **persistent mode:** (not shown in this thesis)


Forkserver

The forkserver makes it possible to avoid calling `execve` for every input, performing it only once at the start of execution. The forkserver is instrumented before the main function of the target application. The AFL++ forkserver protocol can be summarized as follows: ²

1. The forkserver sends the “Welcome” message to the fuzzer (0x41464c01).
2. The fuzzer replies with the status code. At this point the forkserver must check whether $status \oplus 0xffffffff = 0x41464c01$. If it is equal, execution continues; otherwise, the forkserver stops here.
3. The forkserver sends the execution options to the fuzzer. These can indicate that the target is running in persistent mode, or other optimizations.
4. The forkserver sends 0x41464c01 to the fuzzer to end the handshake.
5. The fuzzer writes a message to the forkserver (its content is discarded); the handshake is now finished. This message is the “go” signal for the forkserver.
6. The forkserver invokes `fork()`. The parent writes the PID of the child to the fuzzer

²The differences between the AFL forkserver protocol and the AFL++ forkserver protocol lie mainly in the messages exchanged between the forkserver and the fuzzer

Figure 1.9: AFL++ Forkserver Protocol



and then invokes `waitpid()`.

The child continues execution of the main function until it ends.

7. When the parent returns from the `waitpid()` function (i.e. the child has terminated), the exit status of the child is sent to the fuzzer, and the forkserver waits for another “go” from the fuzzer.

The communication between the Fuzzer and the Target shown in fig. 1.9 occurs using two unidirectionally pipes that uses the file descriptor:

- *FORKSRV_FD* (*fd 198*): used by the forkserver to read a message sent by the fuzzer.
- *FORKSRV_FD + 1* (*fd 199*): used by the forkserver to write a message to the fuzzer.

1.6. Goals and Challenges

The main goal of this thesis is to develop a tool capable of inline, static edge-coverage instrumentation that preserves binary correctness while achieving better performance with

respect to DBI and comparable to other state-of-the-art instrumentation technique. I have chosen to implement the SBI approach instead of DBI due to its low overhead. To do this, I will use the GRAPE binary rewriting framework, a novel approach for binary rewriting that preserves binary executability while introducing small overheads compared to a trampoline-based approach. This approach involves several challenges, mainly related to preserving the binary's correctness and behavior. Moreover, the tool must be compatible with the already existing AFL++ fuzzer. This means implementing the forking communication protocol inside the target binary without recompiling it. This has created no small number of problems, mainly regarding the preservation of the ELF layout after adding new sections. Compared to the ZAFL approach, the SHM update instructions are placed directly inside the BB, rather than in a function called at each BB. This approach drastically increases the size of the file, causing many instruction repetitions (since the SHM instructions are always the same), but was chosen because it does not cause a context switch at each BB. In fact, in order to call a function, the state of the registers must be saved and restored. Moreover, the stack must be allocated, and the callee's RIP must be saved on the stack alongside the old base pointer of the stack. Performing all these operations at each BB would introduce significant overhead.

2 | Approach

In this chapter I will disclose the general aspects of the approach regarding the insertion of the Edge Tracing instructions inside an already compiled binary.

Four are three main steps to be fulfilled to have a binary that is able to trace Edge Coverage and to be compatible with AFL/AFL++:

1. Insertion of the forkingserver invocation prior the main (section 2.1).
2. Shared Memory address discovery (section 2.2).
3. Search of instrumentation points to trace Edge Coverage.
4. Instrument the previously found points with Edge Coverage tracing instructions.

2.1. Forkserver Insertion

This has been one of the most crucial points of the whole work. First of all I have tried to instrument the forkingserver function directly inline. This approach presents many problems. The biggest difficulty is the fact that a forkingserver needs some library functions to work properly:

- `read/write` to implement the communication protocol with the fuzzer.
- `fork/waitpid` to implement the forkingserver main logic.

These functions could be part of the target binary (if the binary already uses these functions) or not. In the best case, the binary already uses these functions and they are available. In this case, a call to these functions could be instrumented at any time, taking into consideration the type of linking with which the binary has been compiled.

- In case of *Static Linking*, the function instructions are already inside the target application, because the linker inserted them during compilation. So a call to these functions could be inserted by finding the function address and then instrumenting a call instruction to that address.
- In case of *Dynamic Linking*, the function instructions are not inside the target

application when it is launched, but the function addresses are dynamically resolved by the dynamic linker as needed. In this case, two new sections must be used (the `.plt` and `.got` sections, as seen in section 1.3).

If the functions required by the forklserver are part of the target binary, their calls can be easily instrumented. Depending on whether the binary is statically or dynamically linked, this can be done either by finding the address of the function and calling it directly, or by finding the address of its `.plt` entry and calling that, leaving the dynamic linker to resolve the function. The problem arises when the original binary does not include these functions. In this case, it is almost impossible to modify the `.plt` and `.got` sections to include the new functions while preserving the binary’s layout and correctness in the general case. In fact, in some exceptional cases padding may provide enough space in the `.got` section for a new entry, but this cannot be assumed in general. Therefore, this first approach has been completely discarded. Since the forklserver is more complex than a simple shared map update, due to the presence of many checks on environment variables and communication with the fuzzer, the best approach is to inject the compiled forklserver code into the target through *Shared Object Injection*. In order to achieve Shared Object Injection, the `.dynamic` section of the target binary must be modified. In particular:

- A `DT_NEEDED` entry is added with the name of the shared object to be injected.
1
- A `RUNPATH` entry is added containing the path where the actual shared object is saved.

2

The shared object injected into the target binary contains the code of the forklserver, which must be executed before the main of the binary. In order for a function of a shared object injected into a binary to be executed before the main, that function must be declared as a “constructor”. As stated in section 1.3, by declaring a function as a constructor, its address is placed inside the `.init_array` section and executed before the actual start of the executable’s main. When an executable has some dependencies, the dynamic linker executes the `.init` section and the functions contained in the `.init_array` section before actually starting the main of the executable.³ In fig. 2.1b the execution flow


¹In reality the shared object name is a string that is written into the `.dynstr` section, and the `.dynamic` section specifies the offset of the string written into `.dynstr`.

²The injection could also be done with the `LD_PRELOAD` environment variable. This variable tells the dynamic linker to add the specified libraries as dependencies of the executable. The problem is that at each execution the variable must be reset with the name of the shared object. By adding the `DT_NEEDED` dependency instead, the dependency becomes persistent.


³The `.init` section is the legacy version of the `.init_array` section.

of the dynamic linker can be seen as it initializes the dynamic dependencies of a target binary that has the dynamic dependencies shown in fig. 2.1a.

1. The dynamic linker searches for *DT_NEEDED* entries of the target binary. When a dependency is found, it maps the library into a loaded Virtual Memory Area (VMA) of the process associated with the target binary being executed.
2. For each library found previously, the dynamic linker executes each function found in the *.init_array* section of that shared object.
3. When the dynamic linker finishes, the `_start` function is executed, and it in turn executes `__libc_start_main`, which executes all the functions in the *.init_array* of the target binary.



(a) Layout of a target binary with its dynamic dependencies



(b) Dynamic linker initialization of dynamic dependencies

Figure 2.1: An example of ELF execution with dynamic dependencies and a constructor function

Therefore, injecting a SO that has a function declared as constructor into the target binary could resolve the problem of instrumenting the forkserver so that it is executed

before the main.

2.2. Shared Memory Instrumentation

As stated in the previous chapter, the SHM is declared by the fuzzer and the forkserver can know the ID of the SHM thanks to an environment variable set by the fuzzer. The forkserver is injected as a constructor function of a SO for which a `DT_NEEDED` dependency has been added to the target binary. But, as stated before, the dynamic linker executes the constructor of the SO (and therefore the forkserver) before the actual initialization of the binary. Moreover, the SHM address is saved in a variable of the SO and is not, by default, visible inside the target binary. The best approach that allows a variable to be accessible outside a binary is to export the symbol that represents the variable.

As seen in section 1.3, the relocations for external symbols are resolved during the execution of the dynamic linker before the main (or at the first invocation of the function in the case of lazy binding). Since the SO contains the symbol holding the address of the SHM, the idea to let the target binary know the address of the SHM is:

- The SO exports that symbol, declaring it with `__attribute__((visibility("default")))` ⁴
- Add a relocation entry to the target binary.

By doing this, the dynamic linker, at binary loading, will resolve the relocation, finding the address of the variable that stores the address of the SHM inside the SO, and write it into the `.got` of the target binary.

Since, as stated in section 2.1, it is difficult to add a GOT entry in the general case (because it depends on the presence of available space in the GOT), the chosen approach has been to add a new section to the binary where the needed symbols, after relocation, will be written. Now the complete method to make the SHM accessible to the target instrumented binary can be reviewed:

1. Add a new section to the binary. This section will be needed as a GOT in which the resolved symbol values will be written.
2. Add a dynamic symbol to the binary. This is the local symbol that will be referenced by the relocation table. The name of the local symbol must be the same as the symbol exported by the SO that has already been injected as a dependency of the

⁴It is not mandatory to explicitly write the visibility attribute, because each symbol has “default” visibility by default.

binary.

3. Add a dynamic relocation entry. This relocation entry must link the symbol to be relocated with the memory address at which the resolved value will be written (i.e. the section previously created to host the GOT).




Figure 2.2: Shared Memory relocation between the target binary and the Shared Object

According to fig. 1.5, the relocations are resolved before the forkserver invocation (because the forkserver lies inside the constructor of the shared object). The correct value for the SHM address is obtained by the forkserver, but the symbol is resolved beforehand, and therefore the address of the symbol is written inside the GOT before the variable contains the address of the SHM. When the main starts, the forkserver has already been executed, and dereferencing the relocated symbol will result in the correct value initialized by the forkserver. The process can be seen in fig. 2.2. In fig. 2.2a the GOT contains the address `0x0`, because the relocation has not yet been performed (assume this is after the SO loading but before symbol relocation). Just after the relocation of symbols is performed, the GOT will contain the address of the external variable, but that variable will still contain `0x0`, because the SHM has not been attached yet, as shown in fig. 2.2b. After

forkserver execution, the external variable will also contain the address of the SHM, as shown in fig. 2.2c.

2.3. Instrumentation Points Search

At this point the target binary has been modified with:

- A DT_NEEDED entry has been added to *.dynamic*, pointing to a shared object that contains the forkserver declared as a constructor, alongside a RUNPATH that points to that shared object.
- A symbol representing the pointer to the SHM has been inserted into *.dynsym*.
- A new section that will be used as a personalized GOT is added.
- A relocation that links the symbol to be resolved to the new section is added.

Now the IPs must be identified so that they can subsequently be instrumented. First, a distinction must be made between two types of instrumentation:

- An initial instrumentation is needed in order to initialize some variables needed for the SHM update rule.
- All the other instrumentations perform the actual SHM update.

In the first case, the IP must be as close as possible to the `main` function of the binary. Finding the main address is easy for both stripped and non-stripped binaries, so, once the main address is identified, the tool tries to find the first address after the function preamble. This goal is not always achieved, because there is no unified grammar for the function prologue. It depends on how many registers the callee wants to save on the stack, how many variables the specific function needs ... In general, the first IP lies in the first BB of the main function.

All the other IPs are found with a Depth First Search (DFS) starting from the initial main BB, searching among the successor BBs of the current one. Specifically, once a BB is found and has not already been analyzed or discovered before, it is added to a queue. The IPs discovery procedure can be seen in algorithm 2.1.

Two different lists must be used in order to distinguish between found basic blocks and basic blocks to be instrumented. Moreover, if a basic block ends with a call, the next BB it is not instrumented. The `EXTRACTSUCCESSORS` function will find all the successors of a BB, excluding :

- The successors that are outside the “code” region (so avoiding instrumenting *.plt*

Algorithm 2.1 Instrumentation point search with DFS algorithm

```

1: found  $\leftarrow []$  {found is a list of discovered basic blocks. It contains both the basic blocks
   to be instrumented and the basic blocks for which instrumentation is not required}
2: to_instrument  $\leftarrow []$  {this is the list of basic blocks to be instrumented. It is a subset of
   the “found” list}
3: current = main_bb
4: stack  $\leftarrow []$ 
5: successors  $\leftarrow$  EXTRACTSUCCESSORS(current, found)
6: for s  $\in$  successors do
7:   stack.PUSH((current, s))
8: end for
9: while stack is not empty do
10:  current_edge  $\leftarrow$  stack.POP()
11:  prev  $\leftarrow$  current_edge[0], curr  $\leftarrow$  current_edge[1]
12:  if prev.last_instruction._opcode  $\neq$  “call” then
13:    to_instrument.APPEND(curr) {If the previous basic block ends with a call instruction,
     the current basic block is not instrumented}
14:  end if
15:  found.APPEND(curr)
16:  successors  $\leftarrow$  EXTRACTSUCCESSORS(curr, found)
17:  for s  $\in$  successors do
18:    stack.PUSH((curr, s))
19:  end for
20: end while

```

stubs). For “code” region is intended all the instruction of the main and other functions that are callable from the main. So avoid instrument functions in *.fini_array* and other not interesting part of *.text* section.

- The edges that are already been found (and so are contained in “found” list)
- The edges that are temporarily inside the stack data structure

For the moment, it can be said that this function filters the successors of the current basic block, discarding those that it would make no sense to consider. For example, if at a certain point the binary performs a call to a library function (e.g. `printf`), without excluding this case the program would also instrument the *plt* stub of the library function.

In fig. 2.3 a disassembled CFG obtained with Cutter [24], representing a program that computes the power operation, can be seen. With respect to fig. 2.3a, the IPs found are:

- **0x11c2**: initial main BB. This is where the initial instrumentation will be placed.
- **0x11f1** (*prev*=0x11c2, *next*=0x11f1)


- ***0x11db*** (prev=*0x11c2*,next=*0x11db*)
- ***0x1253***
 - (prev=*0x11f1*,next=*0x1253*)
 - (prev=*0x11db*,next=*0x1253*)

Apart from the call to the `my_pow` function, all the other functions are exported from `libc`, and so do not need to be instrumented.


With respect to fig. 2.3b, the IPs found are:

- ***0x11b5***
 - (prev=*0x1189*,next=*0x11b5*)
 - (prev=*0x11a7*,next=*0x11b5*)
- ***0x11a7*** (prev=*0x11b5*,next=*0x11a7*)
- ***0x11bd*** (prev=*0x11b5*,next=*0x11bd*)

The initial BB of the function is not instrumented.



(a) Main CFG



(b) Pow function CFG

Figure 2.3: Example of the CFG of a program that computes the power operation between two numbers

2.4. Actual Instrumentation

In the last section, the method by which the IPs are found was discussed. In this section the actual instrumentation is reviewed. As stated in section 2.3, there are two types of instrumentation:

- The first one is used to set up the variables used by the SHM update rule.
- All the other IPs are used for the actual SHM update.

With reference to algorithm 1.3, there are 3 main variables to consider:

- *current_bb_id*: ID of the BB in which the instrumentation is placed
- *previous_bb_id*: ID of the BB from which the branch is taken
- *shared_map*: Address of the SHM

In the last example of the previous section (fig. 2.3) the IPs have already been found. Without considering the first one (which is where the initial instrumentation is placed), for all the other IPs the initial address of the BB is where the instrumentation will be placed. Let us consider, for example, the **0x11f1** IP, with previous BB equal to *0x11c2* and current BB equal to *0x11f1*. In this case, since there is only one incoming edge to **0x11f1**, the *previous_bb_id* can be placed inline. But in the general case (e.g. **0x11b5**) there are multiple edges that end in the BB. For this reason, it becomes impossible to place the ID of the BB directly inline. The solution implemented in the tool discussed in this thesis consists of storing a variable in a section of the binary.⁵ This variable represents the *previous_bb_id* and is initialized to a random value by the initial instrumentation.

For the *current_bb_id*, instead, there is no need to store it in memory. It will be hardcoded (as a random value) in the assembly instrumentation inserted within the basic block.

The *shared_map* has already been resolved and written in the section added to the binary specifically to store resolved symbols.

So now the whole instrumentation phase can be reviewed in greater detail.

Initial Instrumentation

This instrumentation is needed to:

⁵This section could be the same one used for relocation or a new one added specifically to store the data.

- Initialize the *previous_bb_id*. This value is stored in a specific section created for the purpose of holding temporary values needed for instrumentation.
- Dereference the resolved SHM address. This initialization step is performed for optimization purposes. As can be seen in fig. 2.2, the instrumented section that works as a GOT holds the address of the variable inside the injected SO. In order to obtain the actual SHM address, the value must be dereferenced twice. If this double dereference had to be performed at every IP, it would result in a lot of useless overhead. For this reason, in the initial instrumentation the value is dereferenced twice and then saved in the same section as *previous_bb_id*. By doing this, at each IP the SHM address only has to be loaded into a register, and the offset of the correct cell computed.

SHM update instrumentation

algorithm 2.2 shows the general approach to the SHM update. Since the value of the **rax** register is modified, it must be saved before the actual SHM update and restored afterwards. The value of **rax** is not saved on the stack but in the same section where *previous_bb_id* and *shared_map_address* are stored.

Algorithm 2.2 Pseudo-code of Shared Memory update instrumentation

- 1: Save the state of **rax** register that will be used for computation
 - 2: $\text{rax} \leftarrow \text{previous_bb_id}$
 - 3: $\text{rax} \leftarrow \text{rax} \oplus \text{current_bb_id}$ {**rax** will contain the result of the hash operation between the basic blocks}
 - 4: $\text{rax} \leftarrow \text{rax} + \text{shared_map_address}$ {since the shared map is an array, the offset is obtained by adding the initial address to the index}
 - 5: $[\text{rax}] \leftarrow [\text{rax}] + 1$ {the $[\]$ is the dereference operator}
 - 6: $\text{previous_bb_id} \leftarrow \text{current_bb_id} \gg 1$
 - 7: Restore the value of **rax** register to the previously saved value
-

Now the complete approach can be reviewed:

- Add a dynamic dependency on a SO that contains the forkserver function.
- Add a symbol that will contain the SHM address.
- Add two new sections. The first will be used as a GOT to relocate the symbol. The second will be used as temporary memory space to store *previous_bb_id* and *shared_map_address*, alongside register saving.
- Add a relocation entry in order to resolve the symbol previously added. That symbol is exported from the SO.

- Find the IPs with the algorithm discussed in section 2.3.
- Add the initial instrumentation used to initialize the values of *previous_bb_id* and *shared_map_address* stored in the added section.
- For each previously found IP, add the SHM update instructions as seen in algorithm 2.2.

2.5. Forkserver Optimization

Since fuzzing is a technique that relies heavily on execution speed for each single input, optimizing the instrumentation can save time. In the previously discussed approach there is a major cause of overhead: the point at which the forkserver is executed. In fact, as can be seen in fig. 1.5, since the forkserver is executed by the dynamic linker before the execution of the actual target `_start` function, for each input the target binary must execute many operations before actually reaching the main:

- The `_start` function must be executed.
- The `__libc_start_main` function must be executed, in which the environment is set up for the main execution.
- The `.init_array` functions are scanned and, if present, executed.

The ideal approach is to execute the forkserver at the very beginning of the main function, avoiding the execution of all the preamble functions for each input. To achieve this, the relocation mechanism is used again. In the injected SO, the constructor function no longer contains the forkserver, but is used only to prepare its execution. The workflow can be seen in fig. 2.4. Precisely, the constructor function will:

- set up the SHM, reading the SHM ID from the environment variable and assigning its address to the exported symbol.⁶
- assign to a symbol (which will also be exported) the address of the forkserver function that will be executed.

Two symbols will be added to the target binary instead of one:

- The first is the usual SHM address.
- The second is a pointer to the forkserver function.

⁶In the next chapter the SHM attachment will be discussed in depth. In fact, the SO also includes a *standalone mode* in which there is no need for a SHM.

These two symbols are then resolved by the dynamic linker at binary startup and their values are written to the section that has been added. In the initial instrumentation, the address of the forkserver function is first dereferenced and then executed. When the forkserver exits, the rest of the main function is executed, without re-executing all the binary startup functions for each input.




Figure 2.4: Optimized forkserver invocation

3 | Implementation

In the following chapter, the previously discussed approach will be expanded to include the technical specifications and implementation details used to achieve the described solution. Specifically two different files will be examined:

- **libforkserver.so** : this is the SO that will be injected. During the discussion a focus on the forkserver functioning will be faced, alongside the different forkserver functioning modes, the parameters and so on ...
- **instrument.py** : this is the file that actually perform the binary modification and instrumentation. In this case the details on the instrumentation, together with API specification, will be discussed.

3.1. Forkserver Specification

As already stated, the forkserver lies inside a SO that is injected as a dynamic dependency of the binary. Moreover, as seen in section 2.5, the pointer to the forkserver function is initialized (together with the SHM pointer) in the init function (declared as a constructor and executed before the target binary). During the execution of the “init” function, the shared object accesses two environment variables:

- `__AFL_SHM_ID`: represents the SHM ID set by the fuzzer (as seen in section 1.5).
- `__AFL_PLUSPLUS`: this is used to understand which version of the AFL fuzzer is being used. The tool supports both the AFL and AFL++ communication protocols, so this environment variable is used to choose between the two.

Depending on the value of these two environment variables, 3 different forkserver functions can be used, as can be seen in table 3.1.

- **standalone mode**: this is a mode in which the fuzzer is not present. It allows the binary to be run only once (without the typical forkserver loop) and without any communication protocol. In this case there is no SHM, so the piece of memory that

Table 3.1: Different execution modes of the forkingserver based on environment variable values

<code>__AFL_SHM_ID</code>	<code>__AFL_PLUSPLUS</code>	forkserver execution mode
NULL	NULL	<i>standalone mode</i>
NULL	<i>any value</i>	<i>standalone mode</i>
NOT NULL	NULL	<i>afl++ mode</i>
NOT NULL	0	<i>afl mode</i>
NOT NULL	<i>any other value</i>	<i>afl++ mode</i>

is actually updated is allocated by the forkingserver with a `malloc` call.¹

- **afl mode:** In this case there is an actual SHM and a fuzzer that tracks the edge coverage.
- **afl++ mode:** Same as the previous point, but with a different communication protocol between the forkingserver and the fuzzer.

The SO exports two different symbols:

- `uint8_t *__afl_area_ptr`: pointer to the SHM (or the local memory allocated with `malloc`)
- `void (*__forkserver_start_ptr)(void)`: pointer to the forkingserver function that must be invoked by the main

These symbols are resolved inside the target binary in such a way that the binary has access to their values. Depending on the execution mode, these variables have different values.

- **standalone mode:**

- `__forkserver_start_ptr`: The forkingserver function will return without forking, because the target must be executed without tracing.
- `__afl_area_ptr`: Allocated through `malloc`, because it does not need to be a shared map. In this case the edges are traced by the target, although this tracing is not actually used. However, for compatibility with the traced case, `__afl_area_ptr` is still used.

- **afl mode:**

- `__forkserver_start_ptr`: In this case there is actual communication with the

¹this fact is completely transparent to the target binary, which sees the relocated memory address as if it were a shared memory

fuzzer. The protocol used can be seen in algorithm 3.1.

- `__afl_area_ptr`: During the init function, the SHM address is associated with the pointer through a `shmat` function call, passing as parameter the ID of the SHM obtained from the environment variable.
- **afl++ mode**: This mode works in the same way as the previous one, but uses a different communication protocol. The communication protocol is shown precisely in fig. 1.9.

Algorithm 3.1 AFL forkingserver algorithm

```

WRITE(FORKSRV_FD + 1, 0)
while true do
  READ(FORKSRV_FD)
  child_pid ← FORK()
  if child_pid = 0 then
    CLOSE(FORKSRV_FD) {The child must return from the forkingserver and execute the
    main, so the child process can close the file descriptors}
    CLOSE(FORKSRV_FD + 1)
    return
  else
    WRITE(FORKSRV_FD + 1, child_pid) {Communicate the PID of the child to the
    fuzzer}
    status ← WAITPID(child_pid)
    WRITE(FORKSRV_FD + 1, status)
  end if
end while

```

libforkserver.so also has a destructor function called `void afl_forkserver_cleanup(void)`. A function declared as a destructor is a function with the attribute `__attribute__((destructor))`. A reference to a function declared with the previous attribute is placed inside `.fini_array` and is executed after the main. This function is used to detach the SHM from the current process (in the case of afl/afl++ execution mode), or to free the memory in the case of standalone execution mode.

3.2. Instrumentator Specification

In this case the discussion will vert on the **instrument.py** file. This is the file capable of actually instrumenting the target binary. It uses two libraries:

- **lief**: a python library used mainly to add the dynamic dependency, the new sections and the relocation entries. This library is used for the Binary Modification phase.

- **GRAPE**: a binary rewriting framework used to instrument every IP found. This framework is used in the Instrumentation Phase.

3.2.1. Binary Modification

Before actually instrumenting the binary, the previously mentioned modifications must be made. The principal steps are summarized in fig. 3.1.




Figure 3.1: Binary modification phase in which the needed sections, relocations and dynamic dependencies are added. (The dotted arrows specify which information is needed for the relocation phase)

create .my_got section

The created section has the following properties:

- *section type* = *PROGBITS*. This means that the section data is actually stored in the ELF file. For example, the *.bss* section is declared as *NOBITS* because its data is not stored in the ELF file, but a memory region must be allocated when the program starts.
- *section flags*
 - *ALLOC*: the section must occupy actual space in memory. This memory must be allocated during ELF loading.
 - *WRITE*: this section must be writable by the process.
- *section content*: this section must contain the resolved addresses of the two symbols that are relocated. So the content of this section is composed of 16 bytes (in 64-bit addressing, each resolved symbol occupies 8 bytes). The section content is automatically initialized to 0.

- *section alignment* = 8. Since this section will store 8-byte addresses, the alignment is set to 8.

This section, as stated before, must actually be loaded into memory. But this means that the section must be associated with a segment having the LOAD flag, with the same permissions as the section. **lief**, if specified, automatically adds a new segment for each added section.

create symbols

Two symbols must be created. These two symbols correspond to the two exported symbols of the SO, and so their names must match. Specifically, these symbols are created with:

- *symbol name*: since the symbols' names must match, the two symbol names are
 - `__afl_area_ptr`
 - `__forkserver_start_ptr`
- *symbol type* = *OBJECT*. The symbol represents a variable, not a function.
- *symbol binding* = *GLOBAL*. The symbol can be referenced by the dynamic linker. As stated in section 1.3, during ELF loading, the dynamic linker resolves all the symbols for which there is a relocation entry. A symbol can be exported or imported only if its binding is set to GLOBAL. There are LOCAL symbols that, by definition, are local to the object that defines them and therefore cannot be exported (and, consequently, cannot be imported either).
- *symbol visibility* = *DEFAULT*. The symbol has the default visibility of ELF. There are other types (such as HIDDEN or PROTECTED) that are used for local symbols.
- *symbol value* = 0. The symbol does not have a value. This indicates that this symbol must be resolved by the dynamic linker.
- *symbol size* = 8. The symbol size is 8 bytes.

Add relocation entries

Two relocation entries must be added. The first is used to resolve `__afl_area_ptr`, and the second to resolve `__forkserver_start_ptr`. In order to create a relocation entry, the following must be specified:

- *target address of relocation*: where the resolved value will be written
- *relocation type*: for both symbols, `X86_64_GLOB_DAT` is used (the usual type when dealing with GOT resolution)

- *relocation encoding*: for both symbols, RELA encoding is used, which consists in explicitly specifying the addend
- *relocation purpose*: for both symbols, a DYNAMIC purpose is used, because the symbol must be resolved at runtime by the dynamic linker
- *addend* = 0
- *relocation symbol*: the symbol to be resolved

The difference between the two relocation entries lies in the target address and in the relocation symbol. For `__afl_area_ptr`, the target address is the virtual address of the `.my_got` section.

For `__forkserver_start_ptr`, the target address is the virtual address of the `.my_got` section + 8.

Create `.mydata` section

All the considerations made for the previous section are still valid, apart from the section content. In fact, before actually running a target, AFL checks whether the string “`__AFL_SHM_ID`” is contained somewhere in the ELF file, in order to know if the target has been instrumented. If it is not found, AFL stops.² For this reason the aforementioned string is placed inside this section. The string will later be overwritten by other data, but it will pass the AFL check during the fuzzing setup. According to my implementation, this section will contain different pieces of information at different addresses:

- `.mydata+0x10`: contains the ID of the previous BB
- `.mydata+0x20`: contains the address of the SHM initialized by the initial instrumentation
- `.mydata+0x30`: contains the saved `rax` register, to preserve its value before and after the instrumentation

At the end of the **Binary Modification phase**, the modified binary is saved to the filesystem.

3.2.2. Instrumentation

The instrumentation is done entirely using the GRAPE framework. It takes care of adjusting all the reference shifts caused by the altered `.text` section, as well as the eventual


²This check can actually be avoided by setting the environment variable `AFL_SKIP_BIN_CHECK=1`, as must be done when using ZAFL.

resizing of the section (caused by it exceeding its size due to the added instructions), and many other features. GRAPE first reconstructs the CFG of the program and then performs operations on the obtained CFG. The classes that the developed tool will use are:

- **ADDINSTRUCTIONCHANGE**: used to add an instruction at a specific address of the binary CFG. By default, the instruction is added at the specified address, shifting the instruction that was previously there forward.
- **REMOVEINSTRUCTIONCHANGE**: used to remove an instruction from a specific address of the binary.


These two classes take care of CFG modification (splitting a BB, merging two BBs, ...) but do not take care of RIP-relative instruction modification, together with symbol coherence. An example of an error caused by not taking care of modifying single instructions is shown in fig. 3.2.

³ In fig. 3.2 it can be seen that, after having inserted a `nop` instruction, the `lea` that



Address	Instruction	Comment
[0x000011db]		
0x000011db	lea	rax, qword [rip+0xe22] ; str.Usage: __pow_base__exp ; 0x2004
0x000011e2	mov	rdi, rax ; const char *s
0x000011e5	call	puts ; int puts(const char *s)
0x000011ea	mov	eax, 0xffffffff ; -1
0x000011ef	jmp	0x1253

(a) Example of a Basic Block before the addition of an instruction



Address	Instruction	Comment
[0x000011db]		
0x000011db	nop	
0x000011dc	lea	rax, qword [rip+0xe22] ; data.00002005 ; 0x2005
0x000011e3	mov	rdi, rax ; const char *s
0x000011e6	call	puts ; int puts(const char *s)
0x000011eb	mov	eax, 0xffffffff ; -1
0x000011f0	jmp	0x1254

(b) Same Basic Block of fig. 3.2a with a `nop` instruction added

Figure 3.2: Example of the error caused by not taking care of RIP-relative instruction offsets

was at address 0x11db now shifts to address 0x11dc. For this reason the referenced location is no longer a valid data location. Before the addition, the computation was $0x11db + 0xe22 + 7 = 0x2004 < \text{usage_string_location} >$ ⁴. Now the computation is

³As can be seen, the `jmp` at 0x11ef seems to be shifted, but in reality the `jmp` instruction is always RIP-relative and the disassembler used [24] automatically performs the relative address resolution.

⁴The +7 is due to the instruction size

$0x11dc + 0xe22 + 7 = 0x2005$ <?? >

For these reasons, the last step of binary rewriting (after actually adding all the instructions to the BB) uses two other classes:

- **RIPDATA COHERENCE ANALYSIS**: adjusts the RIP-relative instructions.
- **SYMBOL COHERENCE ANALYSIS**: adjusts symbols. For example, after inserting some instructions, the target address of another instruction may have shifted. The reconstruction of the correct target address is performed by the SymbolCoherence phase.

Since GRAPE takes care of all the shifting and adjusting, the instrumentation algorithm only needs to take care of computing the right offset of the RIP-relative instructions to be instrumented and inserting the instructions.

Since the instructions must be added at a specific address, all the offsets are computed as $final_address - instrumentation_address - instruction_size$. For example, with reference to Listing 3.1, supposing the instrumentation address is $0x4100000c$, `offset_to_previous_saved_ID` can be computed as follows:

$$offset_to_previous_saved_ID = (.mydata.virtual_addr + 0x10) - 0x4100000c - 9^5.$$

In algorithm 3.2 the initial instrumentation algorithm can be seen. As stated before, GRAPE automatically increases the size of the `.text` section if, after adding instructions, the section becomes too small to contain them all. When this happens, the `.mydata` and `.my_got` sections could be shifted forward, and the previously computed offsets for the RIP-relative instructions are no longer valid. For this reason, if a shift is detected, all the instructions are recomputed. Moreover, it can be seen that the instructions are inserted in reverse order (the loop index starts from the last instruction). This is because the instructions are inserted one on top of the other.

Initial Instrumentation

Listing 3.1 performs three main tasks:

- Finds the address of the `forkserver` function by dereferencing the resolved symbol, stores it in `rax`, and then performs an indirect call.
- Writes a random integer (between 0 and 65535^6) as the previous BB ID.
- Dereferences the address of the SHM resolved by the dynamic linker and writes it into the `.mydata` section, in order to avoid dereferencing it every single time.

⁵The -9 at the end is needed because the instruction at line 6 of Listing 3.1 uses 9 bytes

⁶Since the AFL shared map size is 65536, each Basic Block ID must be between 0 and 65535

Algorithm 3.2 Initial Instrumentation algorithm**Require:** *addr, mydata_addr, mygot_addr***Ensure:** all the initial instructions are inserted in the correct order at address “addr”

```

1: instructions  $\leftarrow$  COMPUTEINITIALINSTRUMENTATION(addr, mydata_addr, mygot_addr)
   {COMPUTEINITIALINSTRUMENTATION computes the right offsets and returns a list of as-
   sembly instructions written as strings}
2: for i = LEN(instructions) to 1 do
3:   instr  $\leftarrow$  ASM(instructions[i]) {the ASM function converts the assembly instruction,
   passed as a string parameter, into machine code}
4:   ADDINSTRUCTIONCHANGE(addr, instr).MODIFY(binary.cfg)
5:   if .mydata section shifted after .text enlargement then
6:     mydata_addr  $\leftarrow$  mydata.virtual_address
7:     mygot_addr  $\leftarrow$  mygot.virtual_address
8:     instructions  $\leftarrow$  COMPUTEINITIALINSTRUMENTATION(addr, mydata_addr, mygot_addr)
   {recompute all the instructions with the new offsets. GRAPE will take care of shifting
   the instructions already added}
9:     instr  $\leftarrow$  ASM(instructions[i])
10:    REMOVEINSTRUCTIONCHANGE(addr).MODIFY(binary.cfg) {remove and reinsert the
    last instruction, to make sure that the offset to the .mydata section is correct}
11:    ADDINSTRUCTIONCHANGE(addr, instr).MODIFY(binary.cfg)
12:  end if
13: end for

```

```

1 push rsi
2 push rdi
3 mov rax, QWORD PTR[rip + offset_to_relocated_forkserver_function]
4 mov rax, [rax]
5 call rax
6 pop rdi
7 pop rsi
8 mov WORD PTR[rip + offset_to_previous_saved_ID], random_int
9 lea rax, [rip + offset_to_relocated__afl_area_ptr]
10 mov rax, QWORD PTR [rax]
11 mov rax, QWORD PTR [rax]
12 mov QWORD PTR[rip + offset_to_saved_shared_memory_address], rax

```

Listing 3.1: Initial Instrumentation**Basic Block Instrumentation**

The instrumentation added for each BB can be seen in Listing 3.2.

```

1 mov QWORD PTR[rip + offset_to_saved_rax], rax
2 movzx rax, WORD PTR[rip + offset_to_previous_saved_ID]
3 xor ax, random_current_basic_block_ID
4 add rax, QWORD PTR[rip + offset_to_saved_shared_memory_address]

```

```

5  inc BYTE PTR[rax]
6  mov WORD PTR[rip + offset_to_precedent_bb_saved],
   right_shifted_current_ID
7  mov rax, QWORD PTR[rip + offset_to_saved_rax]

```

Listing 3.2: Basic Block Instrumentation

In Listing 3.2, all the considerations previously made for the offset computation are still valid. `random_current_basic_block_ID` and `right_shifted_current_ID` are computed by the instrumenter and hardcoded into each BB.

For the BB instrumentation, the algorithm must also take care of address shifting during instrumentation. Let's consider the example of fig. 2.3 and the BB addresses found in that section. These addresses were `[0x11a7, 0x11b5, 0x11bd, 0x11db, 0x11f1, 0x1253]`. The instrumentation seen in Listing 3.2 occupies a total of 44 bytes. Supposing the location `0x11bd` has been instrumented, all the following locations must then be shifted by 44 bytes. And this computation must be done for each IP found.

But this approach is quite impractical. GRAPE, in order to overcome the issue of changing addresses, uses *labels*. It associates a unique label to each BB. The BB can be uniquely retrieved using this label. If the BB changes its starting address, the label will not change. So, during the IP search, instead of keeping the addresses of the basic blocks saved, the tool keeps the labels of the BBs to be instrumented saved, and only afterwards retrieves the addresses of these BBs. The final approach can be seen below.

After the instrumentation of each BB according to algorithm 3.3, the tool performs `RIPDATA COHERENCE ANALYSIS` and `SYMBOL COHERENCE ANALYSIS`, adjusting all the shifted references, and overwrites the file that **lief** had previously written.

Finally, the tool must check that the `e_phoff` field of the ELF header coincides with the actual offset of the **Program Header Table**, and thus with the `p_offset` field of the **PT_PHDR** segment.

In some binaries, adding new segments (for `.mydata` and `.my_got`) and adding instructions to other segments can cause the **Program Header Table** to shift, and the `e_phoff` field not to be updated correctly.

The tool first reads the ELF header using the **lief** library, obtaining `e_phoff`, and then cycles through all the binary segments until it finds the **PT_PHDR** segment. When it finds the segment, it checks whether the two fields coincide. If not, it manually rewrites the byte that represents `e_phoff` (the 0x20th byte of the header), adjusting it to the correct value. The “manual” rewriting is performed without using an actual library, assuming

Algorithm 3.3 Instrument Basic Block algorithm

Require: $bbs = [bb_label_1 \dots bb_label_n], mydata_addr, mygot_addr$

Ensure: all the Basic Blocks are instrumented correctly

```

1: for  $bb\_label \in bbs$  do
2:    $bb \leftarrow binary.GETBASICBLOCKFROMLABEL(bb\_label)$  {"binary" is the main object
   used by GRAPE and is a virtual representation of the ELF it is rewriting}
3:    $addr \leftarrow bb.initial\_address$ 
4:    $instructions \leftarrow COMPUTEBASICBLOCKINSTRUMENTATION(addr, mydata\_addr, mygot\_addr)$ 
5:   for  $i = LEN(instructions)$  to 1 do
6:      $instr \leftarrow ASM(instructions[i])$ 
7:      $ADDINSTRUCTIONCHANGE(addr, instr).MODIFY(binary.cfg)$ 
8:     if .mydata section shifted after .text enlargement then
9:        $mydata\_addr \leftarrow mydata.virtual\_address$ 
10:       $mygot\_addr \leftarrow mygot.virtual\_address$ 
11:       $instructions \leftarrow COMPUTEBASICBLOCKINSTRUMENTATION(addr, mydata\_addr, mygot\_addr)$ 

12:      $REMOVEINSTRUCTIONCHANGE(addr).MODIFY(binary.cfg)$ 
13:      $ADDINSTRUCTIONCHANGE(addr, instr).MODIFY(binary.cfg)$ 
14:   end if
15: end for
16: end for

```

little-endian addressing. This is done in order to avoid libraries (such as **lief**) adjusting other fields.

4 | Experiments

This chapter presents the objectives of this thesis, the dataset used to verify whether these objectives are achieved, the experimental setup in which the tests are executed, and the tests performed, together with their results.

4.1. Goals

In order to prove that the instrumentation approach is valid, I have chosen three main goals:

- **Goal I** : The instrumented binary must be compatible with the original one. This means that the behavior of the program must be preserved after the instrumentation.
- **Goal II** : The instrumented binary must be compatible with the AFL and AFL++ communication protocol.
- **Goal III** : Show that the instrumentation overhead is competitive with other binary-only rewriters and lower than that of dynamic binary translators.

In order to prove that the tool satisfies these goals, the following experiments were conducted:

- **Experiment I** : Test the original binary and the instrumented binary with the same input, without using a fuzzer, in order to prove that the behavior of the program is preserved after the instrumentation (Goal I).
- **Experiment II** : Perform a fuzzing campaign with three different instrumentations:
 - AFL++ in QEMU mode
 - ZAFI
 - Mine

The chosen fuzzer is AFL++. This test aims to prove Goal II and Goal III.

4.2. Datasets and Experiment Setup

Experiment I

The experiment consists in feeding the same input to both the original binary and the modified one, and checking whether the outputs coincide. The dataset is taken from the *real-world software* dataset used in the Z AFL paper [22]. Not all the software was used. Table 4.1 lists the binaries used, with their versions. All the binaries considered for this

Table 4.1: Binaries extracted from Z AFL *real-world software* dataset

Binary name	Version
bsdtar	bsdtar 3.7.2 - libarchive 3.7.2
readelf	GNU Binutils 2.42
sfconvert	0.3.6
tcpdump	4.99.4
unrtf	0.21.10

experiment are stripped.

Experiment II

The dataset used for the second experiment is LAVA-M [9], in the same version used by Z AFL [22]. This dataset is composed of four coreutils binaries (*base64* , *md5sum* , *uniq* , *who*), to which synthetic bugs were added with the LAVA framework. Each bug has a unique ID and is protected by a checksum. An input can trigger a synthetic bug only if random mutation manages to satisfy the checksum. When a specific bug is triggered, the unique ID of that bug is printed and the program then crashes intentionally.

The coreutils version used is **8.24**. As reported by Chen et al. [7], the LAVA-M version of *md5sum* was modified incorrectly: when run with the *-c* option, it crashes on every input. This affects our second experiment, in which the fuzzing seed must be parsed with *-c*. Since every input triggers a crash, the results for this binary are not meaningful. The experiment was run with AFL++ version 5.03c, over a 24-hour fuzzing campaign with three repetitions for each tested binary. The experiment was executed inside a KVM virtual machine (full virtualization) exposing 16 virtual CPUs, hosting Ubuntu 24 as guest operating system, with 32 GB of RAM. All the different runs were executed concurrently; each was pinned to a dedicated virtual CPU. The 3 different instrumentation methods are tested independently; parallelly are executed the 3 runs of one of the instrumentation method on each of the 4 binaries, having a total of 12 virtual cores occupied. The seeds for each run are different. The initial corpus is composed of only 1 test case.

The general AFL++ command used to each fuzzing campaign is:

`afl-fuzz -i <input_dir> -o <findings_dir> -V 86400 -- <binary> ;` with some difference depending on the binary that is being tested and the instrumentation method:

- For the binaries instrumented with ZAFI the environment variable “AFL_SKIP_BIN_CHECK=1” must be set
- For the binaries fuzzed with AFL++ in QEMU mode, the option “-Q” must be added to the command

Moreover the 4 binaries are executed in different ways:

- *base64* : `base64 -d`
- *md5sum* : `md5sum -c @@`
- *uniq* : `uniq`
- *who* : `who @@`

4.3. Experiment I

Among the binaries listed in table 4.1, *tcpdump* and *unrtf* could not be instrumented: the failure occurs during the GRAPE CFG reconstruction (specifically during Angr CFG analysis).

Each binary is executed different times, with different inputs and different execution options. The output of the original and the instrumented are then compared to check if are the same or not. The results of the experiment are shown in table 4.2.

Table 4.2: Results of experiment I

	Instrumented	total tested inputs	pass	fail
bsdtar	✓	34	34	0
readelf	✓	54	4	50
sfconvert	✓	52	9	43
tcpdump	✗	–	–	–
unrtf	✗	–	–	–

With regard to *readelf*, the instrumented binary encounters a **SIGSEGV** error in all the failed tests. This is due to the fact that, when new sections and segments are added to the binary, the **lief** library automatically reorganizes the ELF layout, causing some sections (e.g. *.bss*) to be shifted. The instructions that referenced these sections will, after the layout change, reference wrong addresses, causing a **SIGSEGV** when they try to

access them.

With regard to *sfconvert* instead, a corruption of the Data Format recognizer had happened. Specifically the instrumented binary could not recognize the Data Format of the audio file anymore after the instrumentation.

A possible reason for the difficulty encountered in instrumenting this dataset could be the fact that these binaries are stripped. The absence of a reason for the impossibility to instrument *tcpdump* and *unrtf* could be the fact that these binaries are stripped. This fact could create difficulties to the CFG reconstruction phase.

4.4. Experiment II

Among the four LAVA-M binaries, none fails during the instrumentation phase. However, two of them (i.e. *md5sum* and *who*) result in an incorrect instrumentation: *md5sum* was already corrupted, as stated before, while *who* was incorrectly instrumented by the tool. The results of the experiment are shown in table 4.3. The runs marked as not done were not executed: since the tool discussed in this thesis produced no valid result for these binaries, there was nothing to compare against. Table 4.3 is divided into four subtables, one per fuzzed binary. Each table reports some parameters obtained for each of the three runs performed.

As can be seen, the tool developed in this thesis finds a considerably smaller number of edges than the other tested methods. This could be caused by a different edge instrumentation method (for example, because the tool does not instrument function calls). Moreover, ZAFL is the only tool capable of executing *md5sum* with the *-c* option but the high number of crashes is probably given by errors of the binary. In fact none of the found unique crashes actually trigger a LAVA Bug, but trigger SIGSEGV errors that make the entire statistic of *md5sum* unreliable.

Table 4.4 shows the size of the instrumented files in MBytes and the average executions per second obtained from table 4.3. As already discussed, the tool developed in this thesis produces larger binaries, because the SHM update is instrumented inline. The average execution speed is only slightly higher than that of ZAFL. Since the tool execution speed, as can be seen in table 4.3, is not always higher with respect to the ZAFL execution speed, and only slightly higher in average, it is not possible to conclude an actual improvement in the execution speed. Worth mentioning that ZAFL performs an optimization of the Intermediate Representation, that surely speeds up the instrumented binary execution. With inline instrumentation it is possible to reach performance that are comparable to the one obtained by ZAFL, without an actual optimization.

(a) base64

	afl			mine			zafl		
	run1	run2	run3	run1	run2	run3	run1	run2	run3
execs_per_sec	398.49	444.98	351.77	1703.45	1738.99	1842.27	1577.45	1792.35	1749.11
execs_done	34429329	38446221	30392855	147177902	150249117	159172058	136291378	154859292	151123034
edges_found	690	690	690	358	370	358	640	672	672
saved_crashes	7	3	2	12	41	11	10	115	118

(b) md5sum

	afl			mine			zafl		
	run1	run2	run3	run1	run2	run3	run1	run2	run3
execs_per_sec	✗	✗	✗	✗	✗	✗	515.09	—	—
execs_done	✗	✗	✗	✗	✗	✗	44503697	—	—
edges_found	✗	✗	✗	✗	✗	✗	9	—	—
saved_crashes	✗	✗	✗	✗	✗	✗	846	—	—

(c) uniq

	afl			mine			zafl		
	run1	run2	run3	run1	run2	run3	run1	run2	run3
execs_per_sec	567.81	576.57	577.17	1758.61	1399.39	1864.43	1773.92	1559.98	1605.66
execs_done	49058377	49815928	49867759	151943556	120907581	161086939	153267103	134782333	138729055
edges_found	473	473	473	242	242	242	443	443	443
saved_crashes	2	0	0	0	0	0	0	1	0

(d) who

	afl			mine			zafl		
	run1	run2	run3	run1	run2	run3	run1	run2	run3
execs_per_sec	225.11	—	—	✗	✗	✗	456.42	—	—
execs_done	19449595	—	—	✗	✗	✗	39434809	—	—
edges_found	13770	—	—	✗	✗	✗	14044	—	—
saved_crashes	0	—	—	✗	✗	✗	2	—	—

Table 4.3: Fuzzing statistics over a 24-hour fuzzing campaign on the LAVA-M dataset using AFL++ in QEMU mode, ZAFL instrumentation + AFL++ fuzzer, and my instrumentation + AFL++ fuzzer

		file_size[MByte]	avg execs/sec
base64	mine	8,25228	1761.57
	zafl	0,137459	1706.3
md5sum	mine	8,256376	✗
	zafl	0,145651	515.09
uniq	mine	8,256376	1674.14
	zafl	0,149747	1646.52
who	mine	8,62092	✗
	zafl	1,718515	456.42

Table 4.4: Instrumented binaries statistics

5 | Limitations and Future Works

As demonstrated by the tests performed and discussed in the previous chapter, the current implementation has some limitations. In this chapter, I will discuss the main limitations of the current implementation and try to outline a direction for future work in this field.

5.1. Limitations

The main limitation is that the developed tool fails to correctly modify some binaries, as shown in table 4.2. I have identified some possible causes of instrumentation failure. One of these occurs in some binaries, where the addition of segments (needed to hold the new sections) causes other sections in other segments to shift. In fact, according to the **LIEF** documentation, LIEF does not guarantee preservation of the ELF layout. For instance, when the dynamic dependency is added to the binary, some sections are modified. Specifically:

- *.dynstr*: the string 'libforkserver.so' must be added to the dynamic string section.
- *.dynamic*: LIEF adds a DT_NEEDED entry that points to the newly added string in *.dynstr*.

If, during either of these two additions, LIEF exceeds the available space in the segment containing the grown section, it reorganizes the segments. This causes a complete displacement of the data sections (such as *.bss*, *.data*, *.rodata*, ...) without a corresponding fix to the instructions that reference these data sections. This can cause a segmentation fault when, for example, an instruction tries to access a value stored in *.bss*. In both PIE and non-PIE binaries, the non-linear displacement between the *.bss* section's address before and after the binary modification causes instructions to no longer point to the correct address. This fact can be seen as the underlying problem both for *readelf* and *sfconvert* of table 4.2. In *who* of table 4.3 instead, the problem needs to be traced in some error inside the instrumentation that cause the disalignment of the instructions of a BB, causing some of these instruction to be illegal instructions.

Moreover, for other binaries, the problem arise in the actual instrumentation phase, high-

lighting a certain fragility of the tool to instrument any type of binary.

Another important limitation that must be mentioned is that the presented approach is not thread-safe. Infact, the `rax` value is saved in a global memory area that is not thread safe. For this reason, if two parallel threads reach a SHM update block at the same time, the value inside of `save_rax` register will be overwritten, causing one of the two threads to restore an incorrect value of `rax`. The same problem happens with previous BB ID saved. Supposing that, at a certain point, four different threads are active, each thread will have, in general, a different previous BB ID.

Another limitation worth mentioning is that the developed tool works only with ELF files, lacking compatibility with the Windows PE executable format.

5.2. Future Works

The main work to be done in the future consists in making this type of solution more resilient to binary corruption. The approach of instrumenting the SHM update directly inline within each BB has proven effective, given the speedup resulting from avoiding a context switch at each BB (which would otherwise be caused by invoking a function) [Reference to the experiments section]. For this reason, a possible direction for future work could be to apply current state-of-the-art techniques to this approach. For example, using the SHM update rule presented in this thesis, several hash collisions occur in the SHM, causing the fuzzer to misinterpret coverage, as demonstrated by Shuitao Gan et al. [12]. In fact, that paper presents a new approach to SHM updates, which was later added as a supported instrumentation method in the AFL++ framework. Implementing a collision-free coverage algorithm as SHM update rule could be a possible improvement over the proposed approach.

Furthermore, implementing a sanitization mechanism could be another research direction building on the work done in this thesis, for example by implementing an Address Sanitizer [26] or a Thread Sanitizer [25].

Bibliography

- [1] AFL++ contributors. AFL++ Frida Mode. https://github.com/AFLplusplus/AFLplusplus/tree/stable/frida_mode, 2026. Accessed: 2026-09-13.
- [2] AFL++ Developers. Fuzzing binary-only targets. URL https://aflplusplus/docs/fuzzing_binary-only_targets/.
- [3] D. Andriesse, X. Chen, V. van der Veen, A. Slowinska, and H. Bos. An in-depth analysis of disassembly on full-scale x86/x64 binaries. In *25th USENIX Security Symposium (USENIX Security 16)*, pages 583–600, Austin, TX, 2016. USENIX Association. ISBN 978-1-931971-32-4. URL <https://www.usenix.org/conference/usenixsecurity16/technical-sessions/presentation/andriesse>.
- [4] D. Bruening, T. Garnett, and S. Amarasinghe. An infrastructure for adaptive dynamic optimization. In *International Symposium on Code Generation and Optimization, 2003. CGO 2003.*, pages 265–275, 2003. doi: 10.1109/CGO.2003.1191551.
- [5] B. Buck and J. K. Hollingsworth. An api for runtime code patching. *The International Journal of High Performance Computing Applications*, 14(4):317–329, 2000. doi: 10.1177/109434200001400404. URL <https://doi.org/10.1177/109434200001400404>.
- [6] M. Böhme, V.-T. Pham, and A. Roychoudhury. Coverage-based greybox fuzzing as markov chain. In *Proceedings of the 2016 ACM SIGSAC Conference on Computer and Communications Security*, CCS16, pages 1032–1043. ACM, Oct. 2016. doi: 10.1145/2976749.2978428.
- [7] P. Chen and H. Chen. Angora: Efficient fuzzing by principled search. In *2018 IEEE Symposium on Security and Privacy (SP)*, pages 711–725, 2018. doi: 10.1109/SP.2018.00046.
- [8] S. Dinesh, N. Burow, D. Xu, and M. Payer. Retrowrite: Statically instrumenting cots binaries for fuzzing and sanitization. In *2020 IEEE Symposium on Security and Privacy (SP)*, pages 1497–1511, 2020. doi: 10.1109/SP40000.2020.00009.

- [9] B. Dolan-Gavitt, P. Hulin, E. Kirda, T. Leek, A. Mambretti, W. Robertson, F. Ulrich, and R. Whelan. Lava: Large-scale automated vulnerability addition. In *2016 IEEE Symposium on Security and Privacy (SP)*, pages 110–121, 2016. doi: 10.1109/SP.2016.15.
- [10] A. Fioraldi, D. Maier, H. Eißfeldt, and M. Heuse. AFL++ : Combining incremental steps of fuzzing research. In *14th USENIX Workshop on Offensive Technologies (WOOT 20)*. USENIX Association, Aug. 2020. URL <https://www.usenix.org/conference/woot20/presentation/fioraldi>.
- [11] A. Fioraldi, D. C. Maier, D. Zhang, and D. Balzarotti. Libafl: A framework to build modular and reusable fuzzers. In *Proceedings of the 2022 ACM SIGSAC Conference on Computer and Communications Security*, pages 1051–1065, 2022.
- [12] S. Gan, C. Zhang, X. Qin, X. Tu, K. Li, Z. Pei, and Z. Chen. Collafl: Path sensitive fuzzing. In *2018 IEEE Symposium on Security and Privacy (SP)*, pages 679–696, 2018. doi: 10.1109/SP.2018.00040.
- [13] P. Godefroid, M. Y. Levin, and D. Molnar. Sage: whitebox fuzzing for security testing. *Communications of the ACM*, 55(3):40–44, Mar. 2012. ISSN 1557-7317. doi: 10.1145/2093548.2093564.
- [14] M. Heuse. Afl-dyninst, 2018. URL: <https://github.com/vanhauser-thc/afl-dyninst>.
- [15] M. Heuse. AFL-DynamoRIO. <https://github.com/vanhauser-thc/afl-dynamorio>, 2018. Accessed: 2026-09-13.
- [16] M. Heuse. AFL-PIN. <https://github.com/vanhauser-thc/afl-pin>, 2018. Accessed: 2026-09-13.
- [17] LLVM Project. libfuzzer – a library for coverage-guided fuzz testing. <https://llvm.org/docs/LibFuzzer.html>. Accessed: 2026-09-11.
- [18] C.-K. Luk, R. Cohn, R. Muth, H. Patil, A. Klauser, G. Lowney, S. Wallace, V. J. Reddi, and K. Hazelwood. Pin: building customized program analysis tools with dynamic instrumentation. *SIGPLAN Not.*, 40(6):190–200, June 2005. ISSN 0362-1340. doi: 10.1145/1064978.1065034. URL <https://doi.org/10.1145/1064978.1065034>.
- [19] S. Malliserry and Y.-S. Wu. Demystify the fuzzing methods: A comprehensive survey. *ACM Computing Surveys*, 56(3):1–38, Oct. 2023. ISSN 1557-7341. doi: 10.1145/3623375.

- [20] Marsman1996. Missing some instrumentations when instrumenting binary with AFL. <https://github.com/HexHive/retrowrite/issues/28>, Aug. 2021. GitHub Issue #28, HexHive/retrowrite repository. Accessed: 2026-09-13.
- [21] B. P. Miller, L. Fredriksen, and B. So. An empirical study of the reliability of unix utilities. *Communications of the ACM*, 33(12):32–44, Dec. 1990. ISSN 1557-7317. doi: 10.1145/96267.96279.
- [22] S. Nagy, A. Nguyen-Tuong, J. D. Hiser, J. W. Davidson, and M. Hicks. Breaking through binaries: Compiler-quality instrumentation for better binary-only fuzzing. In *30th USENIX Security Symposium (USENIX Security 21)*, pages 1683–1700. USENIX Association, Aug. 2021. ISBN 978-1-939133-24-3. URL <https://www.usenix.org/conference/usenixsecurity21/presentation/nagy>.
- [23] O. A. V. Ravnås. Frida: Dynamic instrumentation toolkit for developers, reverse-engineers, and security researchers, 2026. URL <https://github.com/frida/frida>.
- [24] RizinOrg. Cutter, 2026. URL <https://github.com/rizinorg/cutter>. Accessed September 27, 2026.
- [25] J. Schilling, A. Wendler, P. Görz, N. Bars, M. Schloegel, and T. Holz. A binary-level thread sanitizer or why sanitizing on the binary level is hard. In *33rd USENIX Security Symposium (USENIX Security 24)*, pages 1903–1920, 2024.
- [26] K. Serebryany, D. Bruening, A. Potapenko, and D. Vyukov. {AddressSanitizer}: A fast address sanity checker. In *2012 USENIX annual technical conference (USENIX ATC 12)*, pages 309–318, 2012.
- [27] M. Wenzl, G. Merzdovnik, J. Ullrich, and E. Weippl. From hack to elaborate technique—a survey on binary rewriting. 52(3), 2019. ISSN 0360-0300. doi: 10.1145/3316415. URL <https://doi.org/10.1145/3316415>.
- [28] J. Xiao, P. Jiang, Z. Zhao, R. Huang, J. Liu, and D. Li. Robust, efficient, and widely available greybox fuzzing for COTS binaries with system call pattern feedback. In *34th USENIX Security Symposium (USENIX Security 25)*, pages 6239–6258, Seattle, WA, Aug. 2025. USENIX Association. ISBN 978-1-939133-52-6. URL <https://www.usenix.org/conference/usenixsecurity25/presentation/xiao-jifan>.
- [29] M. Zalewski. Technical "whitepaper" for afl-fuzz. https://github.com/google/AFL/blob/master/docs/technical_details.txt. Accessed: 2026-09-19.
- [30] M. Zalewski. American Fuzzy Lop. <http://lcamtuf.coredump.cx/afl/>, 2017. Accessed: 2026-09-13.

- [31] X. Zhu, S. Wen, S. Camtepe, and Y. Xiang. Fuzzing: A survey for roadmap. *ACM Computing Surveys*, 54(11s):1–36, Jan. 2022. ISSN 1557-7341. doi: 10.1145/3512345.

List of Acronyms

BB Basic Block. 2, 4, 5, 15–17, 20, 26–28, 30, 40–44, 53, 54

CFG Control Flow Graph. 2, 4, 17, 27, 41, 49, 50

DBI Dynamic Binary Instrumentation. 3, 4, 14–16, 20

DFS Depth First Search. 26

ELF Executable and Linkable Format. 7, 8, 12, 38–40, 44, 49, 53, 54

GOT Global Offset Table. 8, 9, 11, 13, 24–26, 31, 39

IP Instrumentation Point. 5, 6, 26–28, 30–32, 38, 44

IR Intermediate Representation. 6, 16

JIT Just in-time. 15

PLT Procedure Linkage Table. 8, 9, 11

SBI Static Binary Instrumentation. 3, 14–16, 20

SHM Shared Memory. 17, 18, 20, 24–26, 30–32, 35–37, 40, 42, 50, 54

SO Shared Object. 8, 9, 11, 12, 23–25, 31, 32, 35, 36, 39

TB Translation Block. 15

VMA Virtual Memory Area. 23

List of Figures

1.1	Generic diagram describing the Coverage-Guided Grey-box fuzzing workflow	3
1.2	An example of a Static Binary Rewriting Workflow	5
1.3	An example of Static Instrumentation using the trampoline technique . . .	6
1.4	An example of Lazy Binding involving .plt and .got	10
1.5	Dynamically linked ELF loading procedure	11
1.6	Example of a <i>.dynsym</i> table for a shared object	12
1.7	Example of a <i>.rela.dyn</i> relocation table extracted from a binary	12
1.8	Do-while CFG	17
1.9	AFL++ Forkserver Protocol	19
2.1	An example of ELF execution with dynamic dependencies and a construc- tor function	23
2.2	Shared Memory relocation between the target binary and the Shared Object	25
2.3	Example of the CFG of a program that computes the power operation between two numbers	29
2.4	Optimized forkserver invocation	33
3.1	Binary modification phase in which the needed sections, relocations and dy- namic dependencies are added. (The dotted arrows specify which information is needed for the relocation phase)	38
3.2	Example of the error caused by not taking care of RIP-relative instruction offsets	41

List of Tables

3.1	Different execution modes of the forkingserver based on environment variable values	36
4.1	Binaries extracted from ZAFI <i>real-world software</i> dataset	48
4.2	Results of experiment I	49
4.3	Fuzzing statistics over a 24-hour fuzzing campaign on the LAVA-M dataset using AFL++ in QEMU mode, ZAFI instrumentation + AFL++ fuzzer, and my instrumentation + AFL++ fuzzer	51
4.4	Instrumented binaries statistics	51

List of Algorithms

1.1	Relocation resolution at ELF load time	13
1.2	RESOLVESYMBOL function used by algorithm 1.1	14
1.3	AFL-style SHM update rule	17
1.4	Do-While Example	17
2.1	Instrumentation point search with DFS algorithm	27
2.2	Pseudo-code of Shared Memory update instrumentation	31
3.1	AFL forkserver algorithm	37
3.2	Initial Instrumentation algorithm	43
3.3	Instrument Basic Block algorithm	45

